




Architecture & parameters: $3 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 1$; each hidden layer: Linear $\rightarrow$ LayerNorm $\rightarrow$ SiLU; output: Linear (identity); trainable parameters 10,881 (FC) + 448 (LayerNorm) = 11,329
 Ranges & data: $D \in [0.002, 10.5]$, $z \in [0.2, 10]$, $p \in [0.2, 5000]$, $u \in [0.003, 11.4]$; 64,000 samples; train/validation/test = 80/10/10