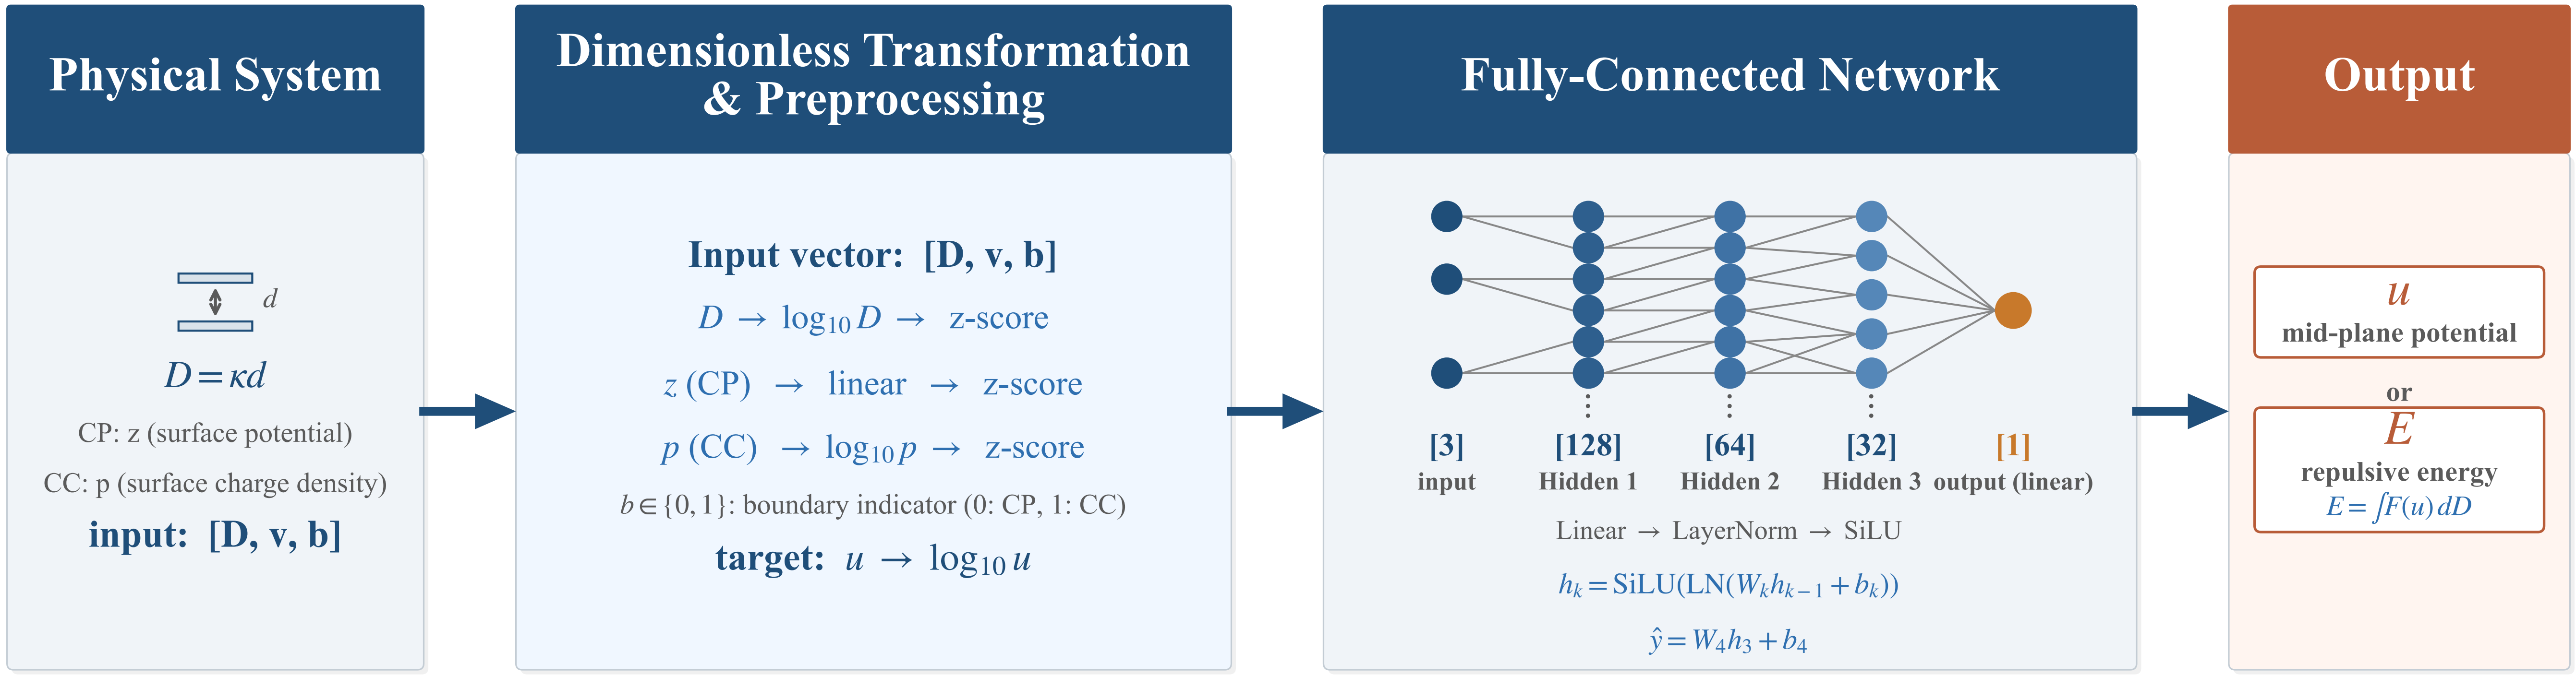

Architecture: 3 → 128 → 64 → 32 → 1; each hidden layer: Linear → LayerNorm → SiLU; output: Linear (identity).

Trainable parameters: 10,881 (FC weights + biases) + 448 (LayerNorm γ, β) = 11,329.

Input ranges: $D \in [0.002, 10.5]$; $z \in [0.2, 10]$; $p \in [0.2, 5000]$; $u \in [0.003, 11.4]$; **dataset: 64,000 samples; train/val/test = 80/10/10.**